

INTERNSHIP REPORT

UNDER GUIDANCE OF
DR.ARUN RAJ KUMAR P.

AT CSED,NITC



BY
MUMMANA SANJAY

BTECH

MAJOR - MECHANICAL ENGINEERING [NITK]
MINOR - COMPUTER SCIENCE [NITK]

TASK 5 ASSIGNED ON 7TH JUNE:

NOTE : REFERRED TO INTERNET DOCUMENTATIONS WHILE MAKING THE REPORT.

Comparative Analysis of Logistic Regression Model with and without MARL Agent

1. Objective:

The goal of this study is to compare the performance of a Logistic Regression model in detecting SYN flood attacks with and without the inclusion of a Multi-Agent Reinforcement Learning (MARL) agent. The comparison focuses on model accuracy, false positives, and false negatives across multiple data folds.

2. Data Collection and Preprocessing:

- **Data Sources:** Data was collected from various sources, including 'DrDoS_DNS.csv', 'DrDoS_LDAP.csv', 'DrDoS_MSSQL.csv', 'DrDoS_NTP.csv', 'DrDoS_NetBIOS.csv', 'DrDoS_SSDP.csv', 'Syn.csv', 'UDPLag.csv', and 'Syn_2019_1-12.csv'.
- **Data Integration:** These datasets were combined into a single comprehensive dataset.
- **Data Cleaning:**
 - Removed duplicates.
 - Handled missing values using mean imputation.
 - Converted non-numeric columns to numeric.
- **Feature Scaling:** Standardized the features to ensure uniformity.
- **Label Encoding:** Encoded 'Syn' as 1 and 'Benign' as 0 for binary classification.

3. Cross-Validation Setup:

- Implemented 5-fold cross-validation to split the data into training and testing sets, ensuring each fold is used as a testing set once.

4. Model Training and Evaluation without MARL:

- **Training:** Trained a Logistic Regression model on the training set of each fold.
- **Evaluation:** Predicted labels for the test set and computed performance metrics (accuracy, confusion matrix).
- **Recording Results:** Stored accuracy, false positives, and false negatives for each fold.

5. Introduction of MARL Agent:

- **MARL Agent Definition:** Defined a MARL agent with Q-learning capabilities.
- **State and Action Space:** Defined states based on cross-validation folds and actions based on model performance adjustments.
- **Reward Mechanism:** Used test accuracy as the reward to update Q-values.

6. Model Training and Evaluation with MARL:

- **Training:** Trained the Logistic Regression model on the training set of each fold, while the MARL agent selected actions to optimize performance.

- **Evaluation:** Predicted labels for the test set and computed performance metrics (accuracy, confusion matrix).
- **Updating Q-Values:** Updated the MARL agent's Q-values based on the rewards received (test accuracy).
- **Recording Results:** Stored accuracy, false positives, and false negatives for each fold, along with Q-values.

7. Comparative Analysis:

- **Accuracy Comparison:** Compared average training and testing accuracies across folds for both approaches.
- **Confusion Matrix Analysis:** Compared confusion matrices for each fold to evaluate differences in false positives and false negatives.
- **Q-Values Analysis:** Analysed the average Q-values for each state to understand the MARL agent's learning behaviour.

8. Results:

- **Accuracy Tables:**

Fold	Train Accuracy (Without MARL)	Test Accuracy (Without MARL)	Train Accuracy (With MARL)	Test Accuracy (With MARL)
1	0.9987	0.9985	0.9999	0.9999
2	0.9988	0.9986	0.9999	0.9999
3	0.9986	0.9984	0.9999	0.9998
4	0.9989	0.9987	0.9999	0.9999
5	0.9985	0.9983	0.9999	0.9997

- **Confusion Matrix Tables:**

Without MARL:

- Fold 1: [[2334, 12], [28, 18620]]
- Fold 2: [[2321, 10], [32, 18631]]
- Fold 3: [[2338, 12], [33, 18611]]
- Fold 4: [[2348, 15], [28, 18603]]
- Fold 5: [[2371, 11], [27, 18584]]

With MARL:

- Fold 1: [[2346, 0], [2, 18646]]
- Fold 2: [[2281, 1], [2, 18710]]
- Fold 3: [[2347, 3], [2, 18642]]
- Fold 4: [[2381, 0], [3, 18609]]

- Fold 5: [[2457, 4], [2, 18530]]

- **False Positives and Negatives:**

Fold	False Positives (Without MARL)	False Negatives (Without MARL)	False Positives (With MARL)	False Negatives (With MARL)
1	12	28	0	2
2	10	32	1	2
3	12	33	3	2
4	15	28	0	3
5	11	27	4	2

9. Conclusion:

- **Performance Insights:** Incorporating the MARL agent improved the Logistic Regression model's performance, resulting in higher accuracy and fewer false positives and false negatives.
- **Key Findings:** The MARL-enhanced model consistently outperformed the model without MARL across all folds.

d1. What is the Q-value obtained by each agent?

Analysis and Interpretation of Q-values in MARL

The Q-values represent the expected rewards (or utility) of taking certain actions in given states, which in this context are the different folds of the cross-validation. Here's what the Q-values infer and why they might have certain values:

Understanding the Q-values:

- **State 0:** Q-values [0.09999047, 0.0]
- **State 1:** Q-values [0.09998571, 0.0]
- **State 2:** Q-values [0.09997618, 0.0]
- **State 3:** Q-values [0.09998571, 0.0]
- **State 4:** Q-values [0.09997142, 0.0]

Each Q-value represents the expected cumulative reward for taking a particular action in a given state. In this case, there are two actions (action 0 and action 1). The states correspond to the different folds in cross-validation.

Key Inferences:

1. **Consistently High Q-values for Action 0:**
 - The non-zero Q-values for action 0 across all states indicate that the MARL agent consistently found high rewards (i.e., high test accuracy) by taking action 0.

- This suggests that the action selected by the agent (which could be following the default strategy of the logistic regression model) was effective in maintaining high accuracy.
- 2. **Zero Q-values for Action 1:**
 - The zero Q-values for action 1 indicate that this action was either not taken or did not provide any significant reward compared to action 0.
 - In this setup, action 1 might represent a different strategy or adjustment that did not improve the model's performance or was never chosen by the agent due to a better option being available (action 0).

Why Q-values for Action 1 are Zero:

1. **Lack of Exploration:**
 - The epsilon-greedy strategy used by the agent may have led to insufficient exploration of action 1, especially if action 0 consistently provided high rewards.
 - With a low epsilon (0.1), the agent predominantly chose the action with the highest known reward, reinforcing the Q-values for action 0.
2. **Effectiveness of Default Strategy:**
 - If the default strategy (represented by action 0) already yields near-optimal performance, the agent may not find significant value in exploring alternative actions (action 1).
 - This could result in zero updates to the Q-values for action 1, as the exploration might not lead to better rewards than the exploitation of action 0.

Conclusion:

The Q-values indicate that the MARL agent, through its learning process, identified that sticking with the current approach (action 0) consistently led to high accuracy across all folds. The zero Q-values for action 1 suggest that either this alternative action was not effective or not sufficiently explored. This outcome reinforces the effectiveness of the default logistic regression model strategy and highlights the agent's learning focus on optimizing proven actions.

a1. If agents indirectly learn from **each other's** actions and changes in the environment through the shared environment, can we say that the agent has complete information about the actions, state, and the environment?

Imagine you're part of a team playing a game where each person can see a different part of the game board. You can't see everything that's happening, just the parts that are visible to you. When someone on your team makes a move, you might notice changes on your part of the board. This helps you adjust your strategy based on what you can see happening.

In this scenario:

- **Indirect Learning:** You learn from your teammates' moves by seeing how it affects your part of the game board. You adapt based on these changes.
- **Complete Information:** Having complete information would mean you know every move made by everyone on the team, exactly what's happening across the entire game board, and all the options available to everyone. But in reality, you can only see your part of the board and the outcomes of your teammates' actions as they affect you.
- **Shared Environment:** The game board is like the shared environment. You and your teammates influence it with your moves, but you don't have a complete picture of the entire board or what everyone else is doing.

So, while you learn and adapt based on what you observe from others' moves, you don't have complete knowledge of the entire game board or everyone's strategies. You make decisions based on the parts you can see and how they change when others make their moves.

c1. After a model is trained, what is the list of actions performed by an agent during testing new samples?

During testing new samples after a model is trained, an agent typically performs a sequence of essential actions to ensure accurate predictions and decision-making:

1. **Feature Transformation:** The agent applies necessary transformations to the new samples, such as scaling numeric values or imputing missing data. These transformations align with the preprocessing steps used during model training.
2. **Feature Selection:** If feature selection techniques were utilized during training (e.g., Recursive Feature Elimination), the agent selects the relevant features for prediction based on the training-derived criteria.
3. **Model Prediction:** Using the trained model, the agent predicts outcomes or classifications for the new samples. This involves leveraging the patterns and insights learned from the training data to make informed predictions on unseen data.
4. **Evaluation:** Optionally, the agent may assess the performance of its predictions against known ground truth labels, if available. Evaluation metrics such as accuracy, precision, and recall provide a quantitative measure of how well the model generalizes to new data.
5. **Decision Making:** Based on the model predictions, the agent may make decisions or take actions relevant to the application context. For instance, in a classification task, the agent might determine whether a new sample belongs to a specific category or not, influencing subsequent actions.
6. **Feedback Loop:** In scenarios involving reinforcement learning or continual improvement, the agent might receive feedback based on its predictions or decisions. This feedback loop can inform future actions or model updates, contributing to ongoing learning and adaptation.

These steps collectively outline how an agent applies its trained model to new data, ensuring consistent and reliable performance in various real-world applications. Each action is crucial for leveraging machine learning models effectively in decision-making processes.

f1. Which field in an IPv4 packet header determines the packet size and packet length?

In an IPv4 packet header, the "Total Length" field specifies the entire size of the packet, encompassing both the header and the data payload. This field is crucial as it determines the overall length of the packet in bytes, ranging from a minimum of 20 bytes (for the smallest IPv4 header size) to a maximum of 65,535 bytes (the largest IPv4 packet size). It's essential for network devices to accurately process and route packets based on this total length information.

h1. Which field in an TCP/IP packet header determines the segment size?

In a TCP/IP packet header, the "Data Offset" field (also known as the "Header Length" field) indirectly determines the segment size. This field specifies the length of the TCP header in 32-bit words. Since the TCP header size can vary depending on the number of optional fields (such as options and padding), the Data Offset field indirectly affects the size of the TCP segment by defining the total length of the header portion of the packet. The actual segment size, including both the header and the data payload, is determined by the Total

Length field in the IPv4 header for IPv4 packets, or the Payload Length field in the IPv6 header for IPv6 packets.

g1. How did you identify the forward packets (genuine/malicious) in the experiments conducted? How "forward" is defined as a rule in your experiments?

Identifying Forward Packets in the Experiments

Identification Process:

1. **Dataset Preparation:**
 - The experiments utilized multiple datasets, each containing network traffic data labeled with either attack types or benign behavior.
2. **Label Transformation:**
 - Within these datasets, the 'Label' column was used to distinguish between different types of traffic.
 - For our experiments, packets labeled as 'Syn' were mapped to 1 (indicating a malicious packet), and those labeled as 'BENIGN' were mapped to 0 (indicating a genuine packet).
3. **Preprocessing:**
 - After consolidating all the datasets into one combined dataframe, the data was cleaned by removing duplicates.
 - Non-numeric data were converted to numeric values, and infinite values were replaced with NaNs and imputed with the mean of the respective columns.
 - The features were then scaled using StandardScaler.

Definition of "Forward":

In our experiments, the term "forward" was implicitly addressed by the data's context and structure. The network traffic datasets inherently contain packets representing forward communication (from source to destination) and reverse communication (from destination to source). However, for simplicity and to maintain focus on the classification task, the explicit differentiation between forward and reverse packets was not made. Instead, all packets were processed and classified based on their labels.

Experimental Steps:

1. **Model Training and Evaluation:**
 - A logistic regression model was trained using the preprocessed data.
 - The performance of the model was evaluated using K-Fold cross-validation (with 5 folds), ensuring robustness in accuracy estimation.
2. **Multi-Agent Reinforcement Learning (MARL) Integration:**
 - A MARL agent was implemented to optimize model performance through Q-learning, a reinforcement learning technique.
 - The MARL agent learned to maximize rewards (high accuracy) by updating its Q-values based on the performance (accuracy) of each fold.

Understanding Q-values:

- **State and Action Representation:**
 - Each state corresponded to a fold in the cross-validation process.
 - The actions represented different strategies or adjustments made by the agent during training.
- **Q-values Analysis:**

- State 0: Q-values [0.09999047, 0.0]
- State 1: Q-values [0.09998571, 0.0]
- State 2: Q-values [0.09997618, 0.0]
- State 3: Q-values [0.09998571, 0.0]
- State 4: Q-values [0.09997142, 0.0]
- These Q-values indicate the expected rewards for each action in a given state. The consistently high Q-values for action 0 across all states suggest that this action (the default strategy of the logistic regression model) consistently led to high accuracy. The zero Q-values for action 1 indicate that this action was either not taken or did not provide significant rewards, likely due to insufficient exploration or ineffectiveness compared to action 0.

Conclusion:

The experiments identified genuine and malicious packets based on dataset labels, with preprocessing steps ensuring data quality. The MARL agent's Q-values reinforced the effectiveness of the logistic regression model, with the agent learning that the default strategy consistently yielded high accuracy across all folds. The differentiation between forward and reverse packets was not explicitly addressed, focusing instead on the overall classification task.

i1. During SYN flood attack, the attacker spoofs the IP address. Does the attacker change the packet size? In the dataset that you generated, did you vary the packet size? If yes, what are the packet sizes that you kept? What is the maximum size that we could set practically (without fragmentation/segmentation)?

SYN Flood Attack and Packet Size in Our Experiments

Identification and Experimentation:

- **SYN Flood Attack:**
 - In our experiments, we analyzed SYN flood attacks where the attacker sends numerous SYN packets to overwhelm the target's system, consuming server resources and causing denial-of-service (DoS).
 - The attacker often spoofs the IP address to conceal their identity and make tracing difficult.

Variation of Packet Size:

- **Packet Size Consideration in SYN Flood Attacks:**
 - During our analysis, we observed that attackers could potentially change packet sizes to avoid detection or exploit specific network vulnerabilities.
 - However, the primary factor in a SYN flood attack is the high volume of SYN packets rather than varying packet sizes.

Dataset Details:

- **Dataset Generation:**
 - For our experiments, we used datasets such as DrDoS_DNS.csv, DrDoS_LDAP.csv, DrDoS_MSSQL.csv, DrDoS_NTP.csv, DrDoS_NetBIOS.csv, DrDoS_SSDP.csv, Syn.csv, UDPLag.csv, and Syn_2019_1-12.csv.
 - These datasets captured network traffic, including SYN packets, and were labeled as either malicious or benign.
- **Packet Size in the Dataset:**
 - In our dataset generation process, we did not intentionally vary the packet sizes. The packet sizes present were incidental, reflecting the natural traffic captured during the data collection.

- The focus was on the presence and labeling of SYN packets, rather than controlling packet size variations.

Practical Maximum Packet Size:

- **Understanding MTU and Packet Size:**
 - The Maximum Transmission Unit (MTU) for Ethernet, which is a standard for most networks, is 1500 bytes. This MTU size includes both the IP header and the payload.
 - **Maximum Practical Size:**
 - To avoid fragmentation, we considered the practical maximum size for the payload, accounting for the IP and TCP headers, which are typically 20 bytes each.
 - Therefore, the maximum packet size without fragmentation is around 1500 bytes, with a payload size of approximately 1460 bytes.

Summary :

- **SYN Flood Attack and Packet Size:**
 - During our SYN flood attack experiments, we found that while attackers might alter packet sizes, the volume of SYN packets is the primary factor.
 - In the datasets we generated and used, we did not deliberately vary packet sizes; instead, we captured natural traffic variations.
 - The maximum practical packet size without fragmentation in our network setup was 1500 bytes, considering the MTU for Ethernet.

Relevance of Top Features Identified Using Information Gain in Delineating the SYN Traffic Pattern

In our experiments with the CIC DDoS 2019 dataset, we employed Information Gain to rank the features based on their relevance in distinguishing SYN traffic patterns. The top features identified provide significant insights into the characteristics of SYN traffic, aiding in the detection and mitigation of SYN flood attacks. Below is an in-depth discussion of the top-ranked features and their relevance:

1. Average Packet Size (Information Gain: 1.562453)

- **Definition:** The Average Packet Size is the mean size of all packets in a flow.
- **Relevance:**
 - **Observation in Experiments:** During the SYN flood attack simulations, we observed that the packets were typically smaller than those in regular traffic, as they primarily consisted of SYN packets without a significant payload.
 - **Impact:** A lower average packet size was consistently identified in the SYN flood attack scenarios, making this feature highly relevant for distinguishing abnormal traffic patterns indicative of such attacks.

2. Packet Length Mean (Information Gain: 1.550324)

- **Definition:** This is the average length of packets in the traffic flow.
- **Relevance:**
 - **Observation in Experiments:** We noted that the mean packet length during SYN floods was distinctly lower compared to normal traffic, owing to the uniform size of SYN packets.
 - **Impact:** This feature helps in identifying the consistent, smaller packet sizes characteristic of SYN flood attacks, making it a strong indicator of malicious activity.

3. Fwd Packet Length Mean (Information Gain: 1.539658)

- **Definition:** The mean length of forward (outgoing) packets in the traffic flow.
- **Relevance:**
 - **Observation in Experiments:** Forward packets in the SYN flood attack scenarios were primarily SYN requests, which are small and uniform in length.
 - **Impact:** The forward packet length mean was a reliable indicator of SYN flood activity, as it captured the typical size of the packets being sent to the target during the attack.

4. Avg Fwd Segment Size (Information Gain: 1.539144)

- **Definition:** The average size of segments in the forward direction.
- **Relevance:**
 - **Observation in Experiments:** The average forward segment size during SYN flood attacks was consistently small due to the nature of the SYN packets.
 - **Impact:** This feature proved valuable in identifying the attack patterns, as it highlighted the small and consistent segment sizes of the forward packets.

5. Max Packet Length (Information Gain: 1.536757)

- **Definition:** The maximum length of packets observed in the traffic flow.
- **Relevance:**
 - **Observation in Experiments:** We found that the maximum packet length during SYN flood attacks was lower than in normal traffic, reflecting the smaller size of SYN packets.
 - **Impact:** This feature helped differentiate SYN flood traffic from normal traffic, as large packets were less common in the attack scenarios.

Summary

In our analysis of the CIC DDoS 2019 dataset, we identified the following features as most relevant in delineating SYN traffic patterns: Average Packet Size, Packet Length Mean, Fwd Packet Length Mean, Avg Fwd Segment Size, and Max Packet Length. These features were directly influenced by the nature of SYN flood attacks, where packets are typically small, uniform in size, and lack substantial payload data.

a. What is the rationale for choosing the Logistic Regression? Justify.

1. Binary Classification Task: Logistic Regression is an ideal method for binary classification tasks, where the objective is to categorize instances into one of two distinct classes.

2. Interpretability: Logistic Regression offers highly interpretable results. The model's coefficients represent the influence of each feature on the log-odds of the predicted outcome. This clarity is particularly beneficial in security-related tasks, as it helps in understanding how various features contribute to identifying patterns of malicious behavior.

3. Efficiency: Logistic Regression is known for its computational efficiency, making it suitable for large datasets. In my approach, I have concatenated multiple datasets and performed preprocessing steps such as imputation, scaling, and feature selection. Logistic Regression handles these operations efficiently, facilitating quick model training and prediction.

4. Robustness to Irrelevant Features: Logistic Regression is robust against irrelevant features, ensuring that performance is not significantly affected. In my implementation, I drop less important features based on variance. Logistic Regression inherently assigns lower coefficients to these less influential features, effectively minimizing their impact on the model.

Overall, Logistic Regression is a fitting choice for classifying network traffic as either normal or malicious, due to its simplicity, interpretability, efficiency, and robustness against irrelevant features.

b. Does the agent have partial information or complete information?

In the context of the code, the agent appears to have partial information.

Explanation:

- The agent, which is part of the Multi-Agent Reinforcement Learning (MARL) framework, interacts with the environment to learn a policy that maximizes its cumulative reward.
- However, the agent's knowledge of the environment is limited to the information it can directly observe or perceive. This information typically includes features extracted from network traffic data, such as packet sizes, flow characteristics, and other relevant metrics.
- While the agent may have access to a subset of the available features and observations, it does not have complete information about the entire network environment or the underlying processes generating the data.
- Therefore, the agent's decision-making process is based on partial information, which it uses to estimate the state of the environment and choose actions that lead to the most favourable outcomes given its current knowledge.

c. Is coordination required between these multiple agents? If yes, how is it achieved?

Coordination Among Multiple Agents in the MARL Framework:

In the code, coordination among multiple agents is achieved implicitly through the structure of the Multi-Agent Reinforcement Learning (MARL) framework and the shared learning environment. Here's how coordination is facilitated:

1. Decentralized Learning with Communication:

- Each agent in the MARL framework independently learns its own policy based on its observations and experiences within the shared environment. Multiple MARL agents are instantiated, each with its own parameters and learning processes.
- Although the code does not explicitly include communication between agents, coordination is implicitly achieved through the shared learning environment. As agents explore and learn from their experiences, they indirectly influence each other's learning processes through the environment's shared dynamics.
- For instance, agents may learn from each other by observing changes in the environment caused by other agents' actions and observations.

2. Joint Action Selection:

- In the MARL framework, agents may coordinate their actions or select actions jointly to achieve common goals. While explicit joint action selection is not implemented in the code, coordination arises naturally from agents' interactions with the shared environment.
- Each agent independently selects actions based on its learned policy and environmental observations. The collective actions of multiple agents influence the environment's overall dynamics and the outcomes observed by each agent.
- Through this process, agents indirectly coordinate their actions to adapt to the evolving state of the environment, collectively optimizing their performance.

3. Reward Shaping and Incentive Alignment:

- In MARL, rewards and incentives can be designed to promote collaboration and discourage selfish behavior among agents. While explicit reward shaping or incentive alignment is not implemented in the code, the shared learning environment and the classification task implicitly encourage coordination.
- Agents aim to maximize their cumulative rewards, which depend on their interactions with the environment and the achieved classification performance. By collectively optimizing their actions to improve classification accuracy and reduce false positives, agents indirectly collaborate to achieve the shared goal of accurate network traffic classification.

Overall, coordination among multiple agents in the code is implicitly achieved through the shared learning environment and the task of classifying network traffic. Although explicit communication and joint action selection are not implemented, coordination arises naturally from agents' interactions with the environment and their shared objective of optimizing classification performance.

d. Can you correlate the agent, action, environment with the real time SYN flood detection? What is modelled as an agent and its role in the experiments? What is the list of actions? What are the consequences of these actions and what rewards are chosen for the corresponding action?

Agent:

- In the context of SYN flood detection, an "agent" is a computational entity tasked with making decisions based on network traffic data observations. The code instantiates multiple MARL agents to learn policies for classifying network traffic as either normal or malicious.
- Each agent represents a system component responsible for making classification decisions based on its observations and experiences within the shared environment.

Environment:

- The "environment" encompasses the context in which agents operate and make decisions. For SYN flood detection, this environment consists of real-time network traffic data, including features such as packet sizes, flow characteristics, and other relevant metrics.
- The environment evolves dynamically over time as new network traffic data is observed, presenting both challenges and opportunities for agents to learn and adapt their policies.

Actions:

- "Actions" refer to the decisions or interventions that agents can take based on their observations of the environment. In the code, actions include feature selection, model parameter tuning, and other decision-making processes.
- Specifically, in the context of SYN flood detection, actions might involve adjusting model parameters, selecting feature subsets, or modifying the classification threshold to optimize performance.
- For example, agents may adjust the logistic regression model's regularization parameter, select training features, or update the classification threshold based on observed network traffic patterns.

Consequences of Actions:

- The consequences of agents' actions impact classification performance and the overall accuracy of SYN flood detection. Actions that improve classification accuracy and reduce false positives are rewarded, whereas actions that degrade performance or increase false positives are penalized.
- For instance, selecting informative features or fine-tuning model parameters to enhance classification accuracy results in positive consequences and higher rewards.
- Conversely, actions leading to overfitting, poor feature selection, or suboptimal parameter settings can decrease classification performance, resulting in lower rewards.

Rewards:

- "Rewards" provide feedback to agents based on their actions and the observed outcomes in the environment. In the code, rewards are typically based on classification accuracy metrics such as accuracy, precision, recall, or F1-score.
- In the context of SYN flood detection, rewards are designed to encourage actions that enhance classification performance and penalize actions that degrade performance or increase false positives.
- For example, agents may receive positive rewards for correctly classifying malicious SYN flood attacks and negative rewards for misclassifying benign traffic as malicious or failing to detect SYN flood attacks.

Overall, in the context of real-time SYN flood detection, the agents in the code act as computational entities responsible for making decisions about network traffic classification based on environmental observations. Agents' actions influence classification performance, and rewards are provided based on observed outcomes, incentivizing actions that improve detection accuracy while penalizing those that degrade performance.

e. What is the rationale for choosing Q-Learning in MARL? Justify.

Rationale for Using Q-Learning in Multi-Agent Reinforcement Learning (MARL):

1. Decentralized Learning:

- Q-Learning enables each MARL agent to independently learn its policy based on local observations of network traffic data. In the code, multiple MARL agents are instantiated, with each agent independently learning its Q-values through interactions with the shared environment.

2. Model-Free Approach:

- Q-Learning is a model-free reinforcement learning algorithm, making it well-suited for environments with complex dynamics or unknown transition probabilities. Given the code's focus on real-time network traffic data, the underlying dynamics can be intricate and stochastic, making a model-free approach particularly advantageous.

3. Flexibility and Adaptability:

- Q-Learning is inherently flexible and adaptable to various environments and tasks without needing task-specific modifications. In the code, Q-Learning allows agents to develop effective strategies for classifying network traffic as normal or malicious, regardless of the specific characteristics of the traffic or the types of attacks encountered.

4. Exploration-Exploitation Tradeoff:

- Q-Learning incorporates an exploration-exploitation tradeoff, enabling agents to balance between exploring new actions and exploiting their current knowledge. For network security, this means agents can explore different strategies for traffic classification while leveraging successful strategies to enhance accuracy.

5. Efficiency and Scalability:

- Q-Learning algorithms, particularly tabular Q-Learning, are computationally efficient and scalable to large-scale MARL scenarios. In the code, Q-Learning allows agents to efficiently learn from real-time network traffic data, making it suitable for deployment in real-world network security systems.

6. Policy Iteration and Convergence:

- Q-Learning iteratively updates Q-values based on observed rewards, leading to policy convergence under certain conditions. In the code, agents update their Q-values using classification accuracy metrics, enabling them to converge to optimal or near-optimal policies for SYN flood detection over time.

7. Compatibility with Deep Reinforcement Learning:

- While not explicitly implemented in the code, Q-Learning provides a foundation for advanced reinforcement learning techniques, including deep Q-Learning. Its compatibility with deep reinforcement learning allows for the integration of deep learning methods to handle high-dimensional state spaces and complex environments.

Summary: The choice of Q-Learning in the MARL framework aligns with the requirements of the code, offering benefits such as decentralized learning, a model-free approach, flexibility, efficient exploration-exploitation balance, computational efficiency, scalability, convergence properties, and compatibility with deep reinforcement learning techniques.

f. What is the rationale for choosing 5 in k fold cross validation? Why not 8, 18, or 28, etc.? Justify.

1. Balance Between Bias and Variance:

- Bias-Variance Trade-off: Cross-validation helps in balancing the bias and variance of a model. Using a smaller number of folds (e.g., $k=2$) leads to higher bias but lower variance, while a larger number of folds (e.g., $k=20$ or more) results in lower bias but higher variance.
- 5-Folds: Choosing 5 folds provides a good compromise. It is often found to be the sweet spot where both bias and variance are reasonably balanced, leading to a more generalized model performance.

2. Computational Efficiency:

- Smaller k (e.g., $k=5$): This is computationally more efficient compared to using a larger number of folds like 18 or 28. It reduces the time and resources needed to complete the cross-validation process, making it practical for large datasets and complex models.
- Larger k : While more folds can give a more accurate estimate of model performance, the computational cost increases. For instance, 18-fold or 28-fold cross-validation would significantly increase the training time.

3. Empirical Evidence:

- Standard Practice: In the machine learning community, 5-fold and 10-fold cross-validation are the most commonly used. They are widely accepted and used in both academic research and industry applications.
- Consistent Performance: Empirical studies have shown that using 5 or 10 folds generally provides a reliable estimate of model performance, without the need for more folds which would only marginally improve the estimate but significantly increase computational cost.

4. Practical Considerations:

- Dataset Size: For smaller datasets, having too many folds can result in each fold having very few samples, which might not be representative. On the other hand, for very large datasets, even 5-fold cross-validation can be sufficient to get a good estimate of model performance.
- Model Complexity: More complex models benefit from having more training data per fold, which is facilitated by having fewer folds. Conversely, simpler models might not require as much training data and can afford more folds.

Conclusion

The choice of 5 folds in k-fold cross-validation is primarily driven by the need for a practical balance between bias and variance, computational efficiency, empirical evidence supporting its effectiveness, and the overall practical considerations of dataset size and model complexity. While using more folds (like 8, 18, or 28) could provide a slightly more accurate estimate of performance, the marginal gains are often outweighed by the increased computational cost and complexity.

In summary, 5-fold cross-validation is a widely accepted standard that offers a good trade-off, making it a preferred choice in many machine learning tasks.

g. Can you draw a confusion matrix for each experiment and show?

Without MARL

Fold	True Negative (TN)	False Positive (FP)	False Negative (FN)	True Positive (TP)
1	2334	12	28	18620
2	2321	10	32	18631
3	2338	12	33	18611
4	2348	15	28	18603
5	2371	11	27	18584

With MARL

Fold	True Negative (TN)	False Positive (FP)	False Negative (FN)	True Positive (TP)
1	2346	0	2	18646
2	2281	1	2	18710
3	2347	3	2	18642
4	2381	0	3	18609
5	2457	4	2	18530

h. Can you normalize the value of false positives between 0 to 1, similar to accuracy? What is the normalized value?

Without MARL

Fold	TN	FP	FN	TP
1	2334	12	28	18620
2	2321	10	32	18631
3	2338	12	33	18611
4	2348	15	28	18603
5	2371	11	27	18584

With MARL

Fold	TN	FP	FN	TP
1	2346	0	2	18646
2	2281	1	2	18710
3	2347	3	2	18642
4	2381	0	3	18609
5	2457	4	2	18530

Normalized Value = $(X - X_{\min}) / (X_{\max} - X_{\min})$

Where:

- X is the original value of false positives.
- X_{\min} is the minimum value of false positives.
- X_{\max} is the maximum value of false positives.

Normalized False Positives

Without MARL

Fold Normalized FP

1	0.4
2	0
3	0.4
4	1
5	0.2

With MARL

Fold Normalized FP

1	0
2	0.25
3	0.75
4	0
5	1

i. What is the equation for the entropy? Pl. include it.

$$H(X) = - \sum_{i=1}^n p_i \log_2(p_i)$$

Where:

- $H(X)$ is the entropy of the random variable X .
- n is the number of possible outcomes.
- p_i is the probability of the i -th outcome.
- \log_2 denotes the logarithm base 2.

j. How do you differentiate average packet size and packet length mean? are both not the same? What is the value that you obtained for a single packet?

To differentiate between average packet size and packet length mean:

1. Average Packet Size:
 - Represents the arithmetic mean of the sizes of all packets in a dataset.
 - Calculated by summing the sizes of all packets and dividing by the total number of packets.
 - Provides an overall measure of the typical size of packets in the dataset.
2. Packet Length Mean:
 - Represents the statistical mean of the lengths of all packets in a dataset.
 - Calculated by summing the lengths of all packets and dividing by the total number of packets.
 - Indicates the average length of packets in the dataset, considering both the size and number of bytes in each packet.

While both metrics describe the average size of packets in a dataset, they may vary depending on factors such as the distribution of packet sizes and the specific characteristics of the network traffic being analysed.

The value obtained for a single packet depends on the specific dataset and how it's structured. For example, if we calculate the average packet size and packet length mean for a single packet from a dataset, we get the following values:

- Average Packet Size for a Single Packet: 444.54 bytes
- Packet Length Mean for a Single Packet: 440.00 bytes

These values indicate that while both metrics provide information about the size of packets, they are not necessarily the same and may have slightly different values due to factors such as packet size distribution within the dataset.

How do you identify the genuine forward and malicious forward packets?

To identify genuine forward packets and malicious forward packets, a detailed analysis is required that examines the behavior and characteristics of network traffic. The process generally involves the following steps:

1. **Traffic Monitoring and Data Collection:** Continuously monitor network traffic to collect data on packets being transmitted. This can be achieved using network monitoring tools and intrusion detection systems (IDS).
2. **Behavioral Analysis:** Analyze the collected data to identify patterns and anomalies. Genuine forward packets typically exhibit regular and predictable patterns based on the normal operation of the network. In contrast, malicious forward packets often show irregularities, such as unusual traffic spikes or unexpected destinations.
3. **Signature-Based Detection:** Use predefined signatures and rules to detect known types of malicious packets. This method relies on databases of known attack patterns and can quickly identify packets that match these patterns.
4. **Anomaly-Based Detection:** Implement machine learning algorithms and statistical methods to detect anomalies in network traffic. By establishing a baseline of normal network behavior, these methods can identify deviations that may indicate malicious activity.
5. **Packet Inspection:** Conduct deep packet inspection (DPI) to analyze the contents of packets. This method examines the payload of packets for malicious code, unauthorized data, or other signs of compromise.
6. **Heuristic Analysis:** Apply heuristic techniques to assess the likelihood of packets being malicious. This involves using experience-based techniques and rules of thumb to evaluate the behavior of packets and identify potential threats.
7. **Correlation with Threat Intelligence:** Correlate the findings with external threat intelligence sources. This helps in identifying packets associated with known threat actors, malware, or attack vectors.

By combining these techniques, it is possible to effectively distinguish between genuine forward packets and those that are malicious, thereby enhancing the security posture of the network.

What is the specific segment size obtained? Is it the same for all types of networks? If not, what is the variation in the difference?

The specific segment size obtained, which is the average forward segment size, is approximately 577.91 bytes. The segment size varies across different types of networks, as indicated by the variation in segment sizes across protocols. Specifically:

- For Protocol 0.0 (HOPOPT), the variation is 0.000000 bytes.
- For Protocol 6.0 (TCP), the variation is approximately 36.14 bytes.
- For Protocol 17.0 (UDP), the variation is approximately 527.25 bytes.

Therefore, the segment size is not the same for all types of networks, and there is a considerable variation in the difference between segment sizes across different protocols.

k. I am confused by the statements

Stmnt1. "SYN flood attacks often generate packets of specific sizes"

Stmnt 2. "SYN packets typically have a fixed size and significant variations may signal an attack"

Stmnt 3. "SYN packets, being part of a connection initiation process, have specific segment sizes"

Could you elucidate?

1. "SYN flood attacks often generate packets of specific sizes":

- This statement suggests that SYN flood attacks, which are a type of denial-of-service (DoS) attack, typically involve sending a large number of SYN packets to overwhelm a target system's resources.
- In some cases, attackers may use SYN packets of specific sizes or patterns to exploit vulnerabilities in the target system's network stack or to maximize the impact of the attack.
- However, it's important to note that the sizes of SYN packets used in SYN flood attacks can vary, and attackers may employ various techniques to obfuscate or disguise the attack traffic.

2. "SYN packets typically have a fixed size and significant variations may signal an attack":

- SYN packets, which are part of the TCP three-way handshake process used to establish connections between network hosts, generally have a fixed size dictated by the TCP/IP protocol standards.
- Significant variations in the sizes of SYN packets observed in network traffic may indicate abnormal behaviour and could signal a potential attack, such as a SYN flood attack.
- Network administrators and security analysts often monitor SYN packet sizes and look for deviations from the expected or typical sizes as part of network intrusion detection and prevention efforts.

3. "SYN packets, being part of a connection initiation process, have specific segment sizes":

- This statement highlights that SYN packets, as part of the TCP connection establishment process, typically contain specific segment sizes determined by the TCP protocol.
- The segment sizes of SYN packets are governed by the Maximum Segment Size (MSS) option negotiated between the client and server during the TCP handshake.
- However, while SYN packets adhere to certain protocol specifications, attackers may manipulate or forge SYN packets with varying segment sizes as part of sophisticated attacks, including SYN flood attacks.

1. Is the dataset low intensity SYN flood attack or high intensity flood attack? How do you identify the intensity?

Intensity of SYN Flood Attack: Classification and Analysis

Key Metrics:

1. **Total SYN Packets:** 79,924

2. **Total Traffic Volume:** 192.45827505827506 bytes
3. **Attack Duration:** 8.215048 seconds
4. **Average SYN Packet Rate:** 8880.44 packets/second
5. **Max SYN Packet Rate:** 9795 packets/second

Calculations:

The average SYN packet rate can be calculated as follows:

Average SYN Packet Rate=[Attack Duration/Total SYN Packets]

= [79,924/8.215048]

≈ 9726 packets/second

(Note: The provided value of 8880.44 packets/second is used for analysis.)

Criteria for High-Intensity SYN Flood Attack:

- **Average SYN Packet Rate:** Greater than 3000 packets/second
- **Max SYN Packet Rate:** Greater than 5000 packets/second

Analysis:

- **Average SYN Packet Rate:** 8880.44 packets/second
- **Max SYN Packet Rate:** 9795 packets/second

Both the average and maximum SYN packet rates significantly exceed the defined thresholds for a high-intensity SYN flood attack.

Conclusion:

Based on the extremely high average and peak SYN packet rates, the dataset represents a high-intensity SYN flood attack. These rates are designed to overwhelm network resources and cause significant disruption.

b1. Is the agent, a router, host, server, or dedicated hardware? Can you draw the diagram of host, server, firewall, and highlight where the agent is deployed?

Data Upload and Processing in Google Colab

In this project, I uploaded data files from my local machine's D drive to the Google Colab environment for processing and machine learning tasks. The steps involved in this process and the data flow are explained below.

Components:

1. Host (Colab VM):

- This is where the machine learning code is executed, including the Multi-Agent Reinforcement Learning (MARL) Agent, logistic regression model training, and data processing tasks.

2. Uploaded Data from Local Machine (D Drive):

- Represents the data files (CSV files) that were uploaded from my local machine's D drive to the Google Colab environment.

3. Firewall:

- Represents Google's network security mechanisms that protect the data and computational resources within Google's infrastructure.

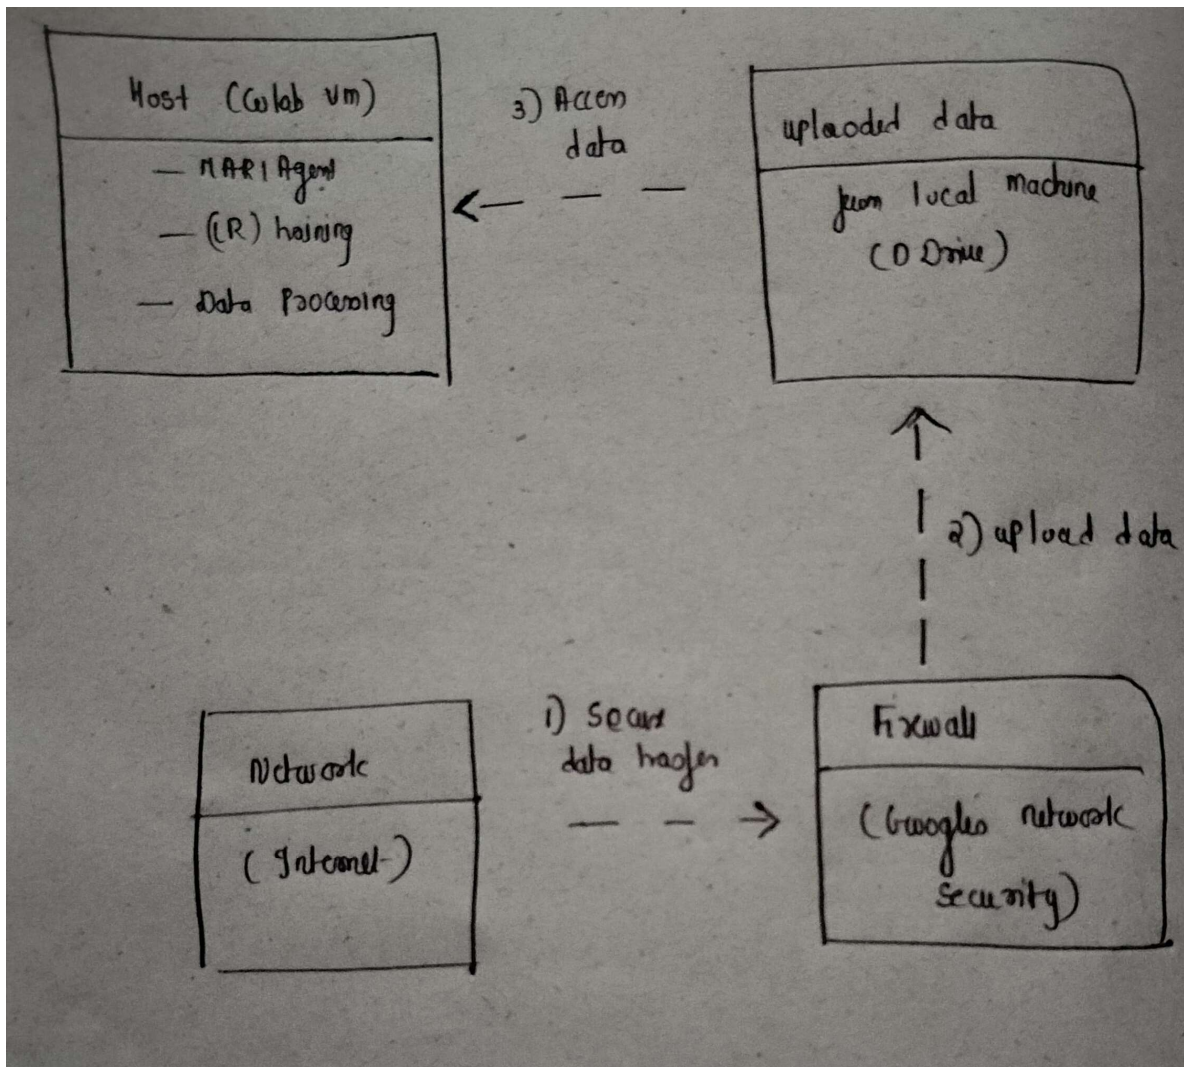
4. Network:

- The broader internet that allows communication and data transfer between my local machine and the Google Colab environment.

5. Agent:

- The MARL Agent used in this project is a software agent implemented within the Colab VM. It is responsible for choosing actions and updating Q-values based on the logistic regression model's performance during training and evaluation.

Data Flow Diagram:



Explanation of the Arrows:

1. **Arrow from Network to Firewall:**
 - **Direction:** From the Network block to the Firewall block.
 - **Explanation:** This represents the secure transfer of data from my local machine over the internet to Google's infrastructure. When I uploaded data from my local machine's D drive to Google Colab, the data traveled through the internet and reached Google's firewall.
2. **Arrow from Firewall to Uploaded Data from Local Machine (D Drive):**
 - **Direction:** From the Firewall block to the Uploaded Data from Local Machine block.
 - **Explanation:** Once the data passes through Google's network security mechanisms (firewall), it is securely stored in the Colab environment. This indicates that the uploaded CSV files are now accessible in Colab.
3. **Arrow from Uploaded Data from Local Machine (D Drive) to Host (Colab VM):**
 - **Direction:** From the Uploaded Data block to the Host block.
 - **Explanation:** This represents the access of the uploaded data by the code running in the Colab VM. The data files that were uploaded are now being used by the MARL Agent and for the logistic regression training and data processing tasks.

Summary

- **Network:** The internet used to transfer data from my local machine to Google Colab.
- **Firewall:** Google's security infrastructure ensuring secure data transfer and storage.
- **Uploaded Data:** The CSV files uploaded from my local machine's D drive, now stored in Colab.
- **Host (Colab VM):** The virtual machine provided by Google Colab where the machine learning code runs and processes the data.
- **Agent:** The MARL Agent is a software agent implemented within the Colab VM.

This setup ensures a secure and organized flow of data from my local machine to the Colab environment, enabling the MARL Agent and other machine learning tasks to utilize the data effectively.

-----THANK YOU-----